



XAI in Medical Image Analysis



Mauricio Reyes, PhD.

Explainable A.I



Source: DARPA

Definitions

- What is Interpretability?

An interpretable machine learning algorithm is described as one where the **link between the features** used by the machine learning and the **prediction itself can be understood by a human**

Another definition

- Produce explanations without sacrificing accuracy
 - Simpler is easier to understand
 - But oversimplified is typically from an accuracy point of view, not interesting




Shortcut learning in deep neural networks

Robert Geirhos^{1,2,4}✉, Jörn-Henrik Jacobsen^{3,4}, Claudio Michaelis^{1,2,4}, Richard Zemel^{3,5}, Wieland Brendel^{1,5}, Matthias Bethge^{1,5} and Felix A. Wichmann^{1,5}

- Principle of “least effort”
- Inductive bias:
 - Model architecture
 - Loss
 - Optimization
 - Training data

Same category for humans
but not for DNNs (intended generalization)

i.i.d.



Domain shift Wang 2018 Adversarial examples Szegedy 2013 Distortions Dodge 2019 Pose Alcorn 2019 Texture Geirhos 2019 Background Beery 2018




o.o.d.



AI for radiographic COVID-19 detection selects shortcuts over signal

Alex J. DeGrave^{1,2,3}✉, Joseph D. Janizek^{1,2,3} and Su-In Lee¹✉



Laterality markers

Attention out of the region of interest

Explaining DL decisions via attention maps




Figure 4: Explaining an image classification prediction made by Google's Inception neural network. The top 3 classes predicted are "Electric Guitar" ($p = 0.32$), "Acoustic guitar" ($p =$





Figure 11: Raw data and explanation of a bad model's prediction in the "Husky vs Wolf" task.



DL learned to detect clinical site!


AI recognition of patient race in medical imaging: a modelling study

Judy Wawira Githaka
Mamungh Chahal

MXR Dense

MXR Dense

- Detect
- Pattern



(0.930-


(0.823-

Interpretability in AI Medical Imaging

Automated classification of Low and High-grade gliomas (LGG vs. HGG)

Q: Are there biases stemming from the data preparation process?

A: Bias of learned patterns detected via interpretability



T1c


T2

Interpretability in AI Medical Imaging

Automated classification of Low and High-grade gliomas (LGG vs. HGG)

Q: Are there biases stemming from the data preparation process?

A: Bias of learned patterns detected via interpretability




T1c

T2

Interpretability in AI Medical Imaging

Automated classification of Low and High-grade gliomas (LGG vs. HGG)

Q: Are there biases stemming from the data preparation process?



T1c

T2


Before

After

Saliency Maps

Interpretability in AI Medical Imaging

Key finding: interpretability enhances data preparation and AI-performance



IMAGENET-TRAINED CNNs ARE BIASED TOWARDS TEXTURE; INCREASING SHAPE BIAS IMPROVES ACCURACY AND ROBUSTNESS

Robert Geirhos
University of Tübingen & IMPRS-IS
robert.geirhos@bethgelab.org

Patricia Rubisch
University of Tübingen & U. of Edinburgh
p.rubisch@sms.ed.ac.uk

Claudio Michaelis
University of Tübingen & IMPRS-IS
claudio.michaelis@bethgelab.org

Matthias Bethge*
University of Tübingen
matthias.bethge@bethgelab.org

Felix A. Wichmann*
University of Tübingen
felix.wichmann@uni-tuebingen.de


Wieland Brendel*
University of Tübingen
wieland.brendel@bethgelab.org




original



greyscale



silhouette



edges



texture

IMAGENET-TRAINED CNNs ARE BIASED TOWARDS TEXTURE; INCREASING SHAPE BIAS IMPROVES ACCURACY AND ROBUSTNESS

Robert Geirhos
University of Tübingen & IMPRS-IS
robert.geirhos@bethgelab.org


Patricia Rubisch
University of Tübingen & U. of Edinburgh
p.rubisch@sms.ed.ac.uk

Claudio Michaelis
University of Tübingen & IMPRS-IS
claudio.michaelis@bethgelab.org

Matthias Bethge*
University of Tübingen
matthias.bethge@bethgelab.org

Felix A. Wichmann*
University of Tübingen
felix.wichmann@uni-tuebingen.de


Wieland Brendel*
University of Tübingen
wieland.brendel@bethgelab.org




original



greyscale



silhouette



edges



texture

IMAGENET-TRAINED CNNs ARE BIASED TOWARDS TEXTURE; INCREASING SHAPE BIAS IMPROVES ACCURACY AND ROBUSTNESS

Robert Geirhos
University of Tübingen & IMPRS-IS
robert.geirhos@bethgelab.org


Patricia Rubisch
University of Tübingen & U. of Edinburgh
p.rubisch@sms.ed.ac.uk

Claudio Michaelis
University of Tübingen & IMPRS-IS
claudio.michaelis@bethgelab.org

Matthias Bethge*
University of Tübingen
matthias.bethge@bethgelab.org

Felix A. Wichmann*
University of Tübingen
felix.wichmann@uni-tuebingen.de

Wieland Brendel*
University of Tübingen
wieland.brendel@bethgelab.org



IMAGENET-TRAINED CNNs ARE BIASED TOWARDS TEXTURE; INCREASING SHAPE BIAS IMPROVES ACCURACY AND ROBUSTNESS

Robert Geirhos
University of Tübingen & IMPRS-IS
robert.geirhos@bethgelab.org


Patricia Rubisch
University of Tübingen & U. of Edinburgh
p.rubisch@sms.ed.ac.uk

Claudio Michaelis
University of Tübingen & IMPRS-IS
claudio.michaelis@bethgelab.org

Matthias Bethge*
University of Tübingen
matthias.bethge@bethgelab.org

Felix A. Wichmann*
University of Tübingen
felix.wichmann@uni-tuebingen.de

Wieland Brendel*
University of Tübingen
wieland.brendel@bethgelab.org



IMAGENET-TRAINED CNNs ARE BIASED TOWARDS TEXTURE; INCREASING SHAPE BIAS IMPROVES ACCURACY AND ROBUSTNESS

Robert Geirhos
University of Tübingen & IMPRS-IS
robert.geirhos@bethgelab.org


Patricia Rubisch
University of Tübingen & U. of Edinburgh
p.rubisch@sms.ed.ac.uk

Claudio Michaelis
University of Tübingen & IMPRS-IS
claudio.michaelis@bethgelab.org

Matthias Bethge*
University of Tübingen
matthias.bethge@bethgelab.org

Felix A. Wichmann*
University of Tübingen
felix.wichmann@uni-tuebingen.de

Wieland Brendel*
University of Tübingen
wieland.brendel@bethgelab.org



Taxonomy

- **Black-boxes** operating methods: No need to have “access” to the internal of models. Also referred to as model-agnostic.
- **White-boxes**: These are methods that require access to the internals of the models. Gradient-based models are one example of white box interpretability models.
- **By output:**
 - Visualization, or saliency maps: Provide pixel-wise values reflecting their importance to the performance of the model being interpreted
 - Concepts: summary statement or keyword.
 - Feature importance: explain models by expressing importance of features E.g. high weights of a model.




Interpretable DL

Convolutional Neural Networks

Typical structure:

- Convolution layers
- Pooling or subsampling layers
- Fully connected layers




Convolutional Neural Networks


Preview

[Zeiler and Fergus 2013]

Visualization of VGG-16 by Lane McIntosh. VGG-16 architecture from [Simonyan and Zisserman 2014].



“Why Should I Trust You?” Explaining the Predictions of Any Classifier: Local Interpretable Model-agnostic Explanations (LIME) – Ribeiro et al . 2017



“Why Should I Trust You?” Explaining the Predictions of Any Classifier: Local Interpretable Model-agnostic Explanations (LIME) – Ribeiro et al . 2017




Figure 4: Explaining an image classification prediction made by Google’s Inception neural network. The top 3 classes predicted are “Electric Guitar” ($p = 0.32$), “Acoustic guitar” ($p = 0.24$) and “Labrador” ($p = 0.21$)





Figure 11: Raw data and explanation of a bad model’s prediction in the “Husky vs Wolf” task.

“Why Should I Trust You?” Explaining the Predictions of Any Classifier: Local Interpretable Model-agnostic Explanations (LIME) – Ribeiro et al . 2017

Local Interpretable Model-Agnostic Explanations (LIME)



Algorithm 1 Sparse Linear Explanations using LIME

Require: Classifier f , Number of samples N

Require: Instance x , and its interpretable version x'

Require: Similarity kernel π_x , Length of explanation K

$\mathcal{Z} \leftarrow \{\}$

for $i \in \{1, 2, 3, \dots, N\}$ **do**

$z'_i \leftarrow \text{sample_around}(x')$


$\mathcal{Z} \leftarrow \mathcal{Z} \cup \langle z'_i, f(z_i), \pi_x(z_i) \rangle$

end for

$w \leftarrow \text{K-Lasso}(\mathcal{Z}, K) \triangleright$ with z'_i as features, $f(z)$ as target


return w

Explaining by finding prototype class members




Simonyan et al. 2013

- Global explanation technique
- Find pattern maximizing class activation
- Not-necessarily a real image








bell pepper




lemon




husky




washing machine



computer keyboard




kit fox



goose



ostrich



limousine



Simonyan et al. 2013. Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps

Synthesizing the preferred inputs for neurons in neural networks via deep generator networks

- Nguyen et al. 2016

Regularization via generative adversarials

- 1.- Forwards pass of gradients
- 2- Detect maximum activation
- 3.- Backward pass of "FC6" into encoder



Feature vector

classifier

generator

$$\hat{\mathbf{y}}^l = \arg \max_{\mathbf{y}^l} (\Phi_h(G_l(\mathbf{y}^l)) - \lambda \|\mathbf{y}^l\|)$$

- Prior generator and DNN can be trained separately

Synthesizing the preferred inputs for neurons in neural networks via deep generator networks

- Nguyen et al. 2016



mosque



lipstick



brambling



leaf beetle



badger



toaster



triumphal arch



cloak



lawn mower



library



cheeseburger



swimming trunks



barn



candle



table lamp



sandbar



French loaf



lemon



chest



running shoe



water jug



pool table



broom



cellphone



aircraft carrier




entertainment ctr



jean

Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization - Selvaraju et al. 2017

- Generalization of CAM: Zhou et al. Learning Deep Features for Discriminative Localization. In *CVPR*, 2016.
 - Applicable to any CNN-based (CAM requires conv feature maps $\rightarrow$ global average pooling $\rightarrow$ softmax layer)
- Motivation – Good visualization is class-specific and detailed



Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization - Selvaraju et al. 2017


Single forward and (partial) backward pass

Neuron importance weights

A^k : kth feature map

y^c : score for class c (before softmax)

$$\alpha_k^c = \overbrace{\frac{1}{Z} \sum_i \sum_j}^{\text{global average pooling}} \underbrace{\frac{\partial y^c}{\partial A_{ij}^k}}_{\text{gradients via backprop}}$$



Source: Wojciech Samek

$$L_{\text{Grad-CAM}}^c = \text{ReLU} \left(\underbrace{\sum_k \alpha_k^c A^k}_{\text{linear combination}} \right)$$

Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization - Selvaraju et al. 2017

Robustness to adversarial examples: Examples where VGG got “fooled”



Boxer: 0.40 Tiger Cat: 0.18

(a) Original image




Airliner: 0.9999

(b) Adversarial image



Boxer: $1.1e-20$

(c) Grad-CAM “Dog”



Tiger Cat: $6.5e-17$

(d) Grad-CAM “Cat”



Interpretable Explanations of Black Boxes by Meaningful Perturbation - Fong et al. 2018

Motivation: Saliency (gradient-based) approaches are not specific enough


Key idea: select the right perturbation* to study the effect on the prediction function $f(x)$

Meaning of an explanation depends on the meaning of the changes applied to the input x


*Perturbation $\leftrightarrow$ type of information deletion



Interpretable Explanations of Black Boxes by Meaningful Perturbation - Fong et al. 2018




Testing with Concept Activation Vectors – TCAV – Kim et al. 2017



Were there more pixels on the cash machine than on the person?

Did the 'human' concept matter?
Did the 'glasses' or 'paper' matter?




Which concept mattered more?

Is this true for all other cash machine predictions?

Wouldn't it be great if we can **quantitatively** measure how important *any* of these **user-chosen concepts** are?

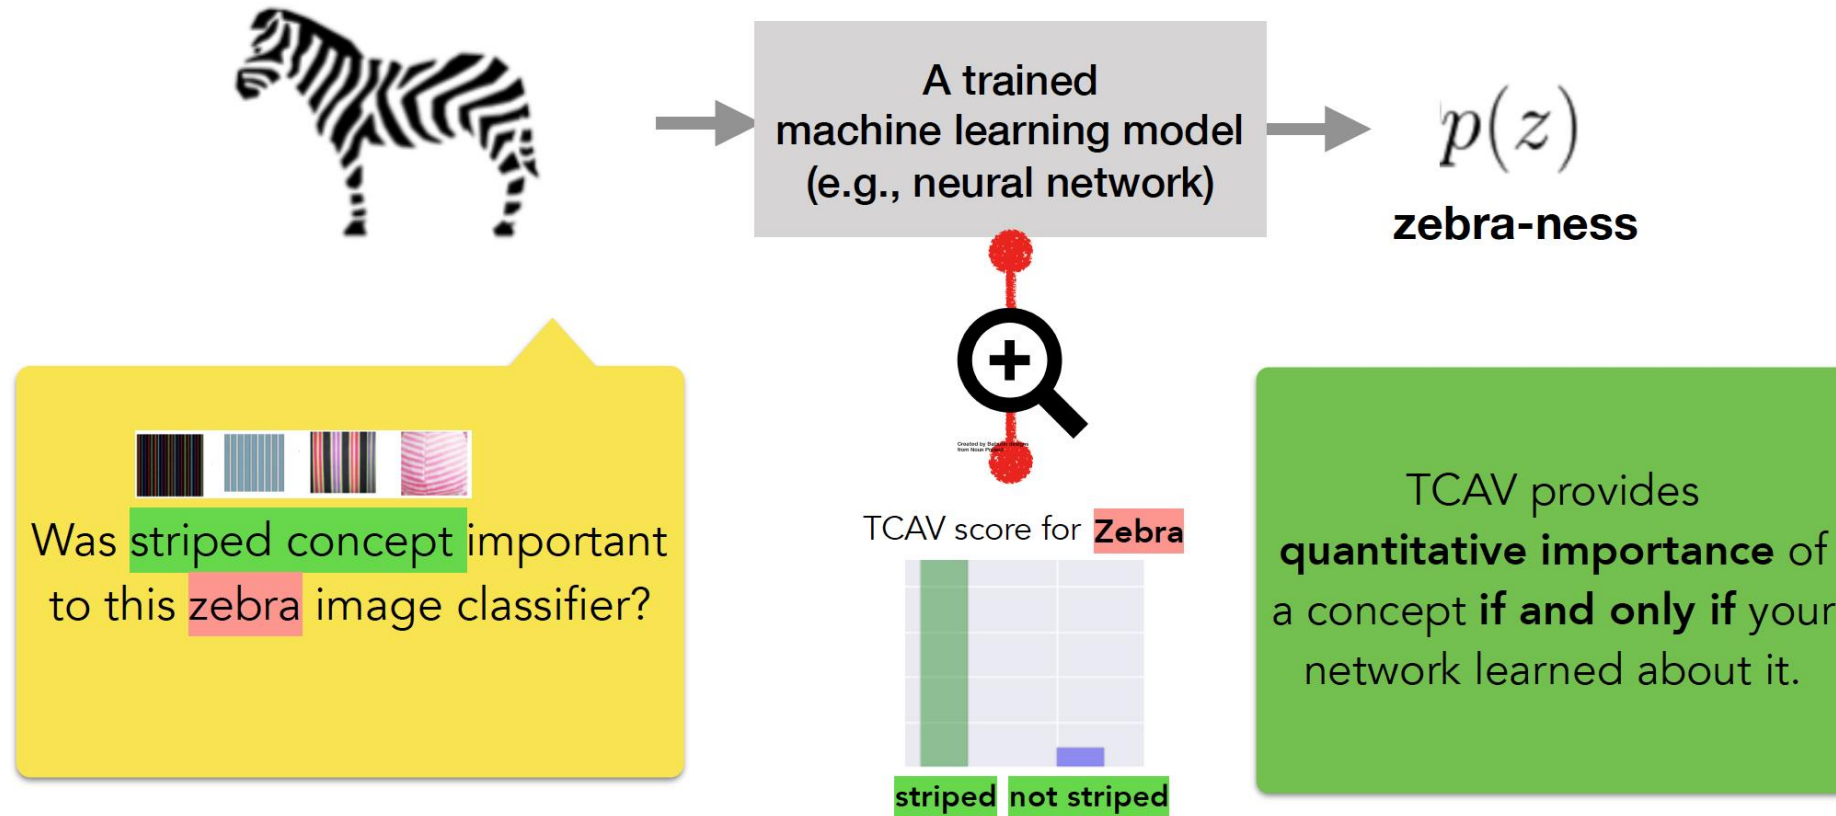
Testing with Concept Activation Vectors – TCAV – Kim et al. 2017




Quantitative explanation: how much a **concept** (e.g., gender, race) was important for a **prediction** in a trained model.

...even if the **concept** was not part of the training.

Testing with Concept Activation Vectors – TCAV – Kim et al. 2017




Testing with Concept Activation Vectors – TCAV – Kim et al. 2017



Testing with Concept Activation Vectors – TCAV – Kim et al. 2017

Inputs:




Train a linear classifier to separate activations.

CAV (v_C^l) is the vector **orthogonal** to the decision boundary.


[Smilkov '17, Bolukbasi '16, Schmidt '15]

Testing with Concept Activation Vectors – TCAV – Kim et al. 2017



Testing with Concept Activation Vectors – TCAV – Kim et al. 2017

TCAV



$$\begin{aligned} \text{zebra-ness} &\rightarrow \frac{\partial p(z)}{\partial \mathbf{v}_C^l} = S_{C,k,l}(\mathbf{x}) \\ \text{striped CAV} &\rightarrow \end{aligned}$$

Directional derivative with CAV





$$\begin{aligned} &S_{C,k,l}(\text{zebra}) \\ &S_{C,k,l}(\text{zebra head}) \\ &S_{C,k,l}(\text{zebra body}) \\ &S_{C,k,l}(\text{zebra tail}) \end{aligned} \quad \left. \vphantom{\begin{aligned} &S_{C,k,l}(\text{zebra}) \\ &S_{C,k,l}(\text{zebra head}) \\ &S_{C,k,l}(\text{zebra body}) \\ &S_{C,k,l}(\text{zebra tail}) \end{aligned}} \right\}$$

$$\text{TCAV}_{Q_{C,k,l}} = \frac{|\{\mathbf{x} \in X_k : S_{C,k,l}(\mathbf{x}) > 0\}|}{|X_k|}$$

fraction of k-class inputs whose l-layer activation vector was positively influenced by concept C


Testing with Concept Activation Vectors – TCAV – Kim et al. 2017

TCAV in Two widely used image prediction models



TCAV facilitates spotting of biases in the datasets

What is AI learning? What if this differs from human reasoning/knowledge?



(a) Husky classified as wolf



(b) Explanation

Shortcut!

VS.

Artificial intelligence predicts patients' race from their medical images

Study shows AI can identify self-reported race from medical images that contain no indications of race detectable by human experts.


Rachel Gordon | MIT CSAIL
May 20, 2022





Knowledge discovery!

Biases!


Right for the Wrong Reasons: Shortcut Learning as a Barrier to Clinical Deployment




Right for the Wrong Reasons: Shortcut Learning as a Barrier to Clinical Deployment




Right for the Wrong Reasons: Shortcut Learning as a Barrier to Clinical Deployment




Right for the Wrong Reasons: Shortcut Learning as a Barrier to Clinical Deployment



Right for the Wrong Reasons: Shortcut Learning as a Barrier to Clinical Deployment






Reliable Clinical AI: Fair and Trustworthy






“Primum non nocere”
(First do no harm)

Reliable Clinical AI: Fair and Trustworthy






*“Primum non nocere”
(First do no harm)*

Reliable Clinical AI: Fair and Trustworthy






“Primum non nocere”
(First do no harm)

Reliable Clinical AI: Fair and Trustworthy





*“Primum non nocere”
(First do no harm)*

Reliable Clinical AI: Fair and Trustworthy





*“Primum non nocere”
(First do no harm)*

Reliable Clinical AI: Fair and Trustworthy




“Primum non nocere”
(First do no harm)

Reliable Clinical AI: Fair and Trustworthy



“Primum non nocere”
(First do no harm)

Reliable Clinical AI: Fair and Trustworthy





Balance between guidance and pattern discovery





Reliable Clinical AI: Fair and Trustworthy

Balance between guidance and pattern discovery




Reliable Clinical AI: Fair and Trustworthy

Balance between guidance and pattern discovery



Inductive
Bias



Reliable Clinical AI: Fair and Trustworthy

Consolidation

Atelectasis




Cardiomegaly

Pleural Effusion

Edema


Reliable Clinical AI: Fair and Trustworthy




Reliable Clinical AI: Fair and Trustworthy



Atelectasis




Pleural Effusion





Reliable Clinical AI: Fair and Trustworthy


Consolidation



Atelectasis




Cardiomegaly




Pleural Effusion




Edema




Reliable Clinical AI: Fair and Trustworthy



Reliable Clinical AI: Fair and Trustworthy




Reliable Clinical AI: Fair and Trustworthy



Reliable Clinical AI: Fair and Trustworthy


Saliency Inductive Bias: SIBNet



Setup:
Five-fold validation
Five classes
DenseNet-121 backbone


Reliable Clinical AI: Fair and Trustworthy

Saliency Inductive Bias: SIBNet



Reliable Clinical AI: Fair and Trustworthy


Saliency Inductive Bias: SIBNet



Reliable Clinical AI: Fair and Trustworthy

Saliency Inductive Bias: SIBNet

Indirect Human-AI Alignment



DenseNet-121
(Baseline)



Pham et al. 2020




SIBNet

Reliable Clinical AI: Fair and Trustworthy

Combining Human-AI alignment and XAI-based Inductive Bias


Reliable Clinical AI: Fair and Trustworthy

Combining Human-AI alignment and XAI-based Inductive Bias






Reliable Clinical AI: Fair and Trustworthy

Combining Human-AI alignment and XAI-based Inductive Bias



Reliable Clinical AI: Fair and Trustworthy

Combining Human-AI alignment and XAI-based Inductive Bias




Zhang et al. Nat. Comm. 2023
KAD – Human-based
Knowledge-Graph guidance




Hybrid
Human-AI & XAI-IB

Saliency maps evolve during training



Key message: Saliency maps can be seen/used as a **fingerprint of model's behavior to inputs**

Graph Node-based Interpretability Guided for Active Learning?




IEEE TRANSACTIONS ON MEDICAL IMAGING, VOL. 42, NO. 3, MARCH 2023


Graph Node Based Interpretability Guided Sample Selection for Active Learning

Dwarikanath Mahapatra^{id}, Alexander Poellinger^{id}, and Mauricio Reyes^{id}


Graph Node-Based Interpretability Guided Sample Selection (GESTALT)



Graph Node-Based Interpretability Guided Sample Selection (GESTALT)





XAI-boosted Active Learning



Mmh, easy sample ...

Mahapatra et al. MedIA 2024
Mahapatra et al. MedIA 2024b
Shu et al. Miccai-iMIMIC 2024
Mahapatra et al. TMI 2022



XAI-boosted Active Learning



Mmh, a hard (a.k.a **informative**) sample...
I can learn from it!

Mahapatra et al. MedIA 2024
Mahapatra et al. MedIA 2024b
Shu et al. Miccai-iMIMIC 2024
Mahapatra et al. TMI 2022



XAI-boosted Active Learning



Mmh, a hard (a.k.a **informative**) sample...
I can learn from it!

Mahapatra et al. MedIA 2024
Mahapatra et al. MedIA 2024b
Shu et al. Miccai-iMIMIC 2024
Mahapatra et al. TMI 2022



XAI-boosted Active Learning



Mmh, a hard (a.k.a **informative**) sample...
I can learn from it!

Mahapatra et al. MedIA 2024
Mahapatra et al. MedIA 2024b
Shu et al. Miccai-iMIMIC 2024
Mahapatra et al. TMI 2022



XAI-boosted Active Learning



Mmh, a hard (a.k.a **informative**) sample...
I can learn from it!

Mahapatra et al. MedIA 2024
Mahapatra et al. MedIA 2024b
Shu et al. Miccai-iMIMIC 2024
Mahapatra et al. TMI 2022



XAI-boosted Active Learning



Mmh, a hard (a.k.a **informative**) sample...
I can learn from it!

Mahapatra et al. MedIA 2024
Mahapatra et al. MedIA 2024b
Shu et al. Miccai-iMIMIC 2024
Mahapatra et al. TMI 2022


XAI-boosted Active Learning



Mmh, a hard (a.k.a **informative**) sample...
I can learn from it!


Mahapatra et al. MedIA 2024
Mahapatra et al. MedIA 2024b
Shu et al. Miccai-iMIMIC 2024
Mahapatra et al. TMI 2022

XAI-boosted Active Learning



Let's wrap-up

- Interpretability of AI systems is propelled by new findings from the DL community to safely translate this technology to (medical) applications.
- The goal of interpretability is to have enough information for the task at hand. In medicine: Patient-safety is first.
- Several types of XAI methods
- Current challenges: XAI needs to be multimodal and longitudinal to match clinical workflows.
- XAI in the era of Foundation Models: Bigger Black Box?



Bibliography

- Adadi, A., & Berrada, M. (2018). Peeking Inside the Black-Box: A Survey on Explainable Artificial Intelligence (XAI). *IEEE Access*, 6, 52138–52160. <http://doi.org/10.1109/ACCESS.2018.2870052>
- Adebayo, J., Gilmer, J., Muelly, M., Goodfellow, I., Hardt, M., & Kim, B. (2018). Sanity Checks for Saliency Maps. Retrieved from <http://arxiv.org/abs/1810.03292>
- Barocas, S., Friedler, S., Hardt, M., Kroll, J., Venka-Tasubramanian, S., & H. Wallach, H. (2018). The FAT-ML Workshop Series on Fairness, Accountability, and Transparency in Machine Learning.
- Caruana, R., Lou, Y., Gehrke, J., Koch, P., Sturm, M., & Elhadad, N. (2015). Intelligible Models for HealthCare. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining - KDD '15* (pp. 1721–1730). New York, New York, USA: ACM Press. <http://doi.org/10.1145/2783258.2788613>
- Datta, A., Sen, S., & Zick, Y. (2016). Algorithmic Transparency via Quantitative Input Influence: Theory and Experiments with Learning Systems. In *2016 IEEE Symposium on Security and Privacy (SP)* (pp. 598–617). IEEE. <http://doi.org/10.1109/SP.2016.42>
- Doshi-Velez, F., & Kim, B. (2017). Towards A Rigorous Science of Interpretable Machine Learning, (MI), 1–13. <http://doi.org/10.1016/j.intell.2013.05.008>
- Eaton-Rosen, Z., Bragman, F., Bisdas, S., Ourselin, S., & Cardoso, M. J. (2018). Towards safe deep learning: accurately quantifying biomarker uncertainty in neural network predictions. Retrieved from <http://arxiv.org/abs/1806.08640>
- Gale, W., Oakden-Rayner, L., Carneiro, G., Bradley, A. P., & Palmer, L. J. (2018). Producing radiologist-quality reports for interpretable artificial intelligence. Retrieved from <http://arxiv.org/abs/1806.00340>
- Gallego-Ortiz, C., & Martel, A. L. (2016). Interpreting extracted rules from ensemble of trees: Application to computer-aided diagnosis of breast MRI. In *ICML Workshop on Human Interpretability in Machine Learning (WHI)*. Retrieved from <http://arxiv.org/abs/1606.08288>
- Gillies, R. J., Kinahan, P. E., & Hricak, H. (2016). Radiomics: Images Are More than Pictures, They Are Data. *Radiology*, 278(2), 563–77. <http://doi.org/10.1148/radiol.2015151169>
- Gilpin, L. H., Bau, D., Yuan, B. Z., Bajwa, A., Specter, M., & Kagal, L. (2019). *Explaining Explanations: An Overview of Interpretability of Machine Learning*. Retrieved from <https://arxiv.org/pdf/1806.00069.pdf>
- Goldstein, A., Kapelner, A., Bleich, J., & Pitkin, E. (2015). Peeking Inside the Black Box: Visualizing Statistical Learning With Plots of Individual Conditional Expectation. *Journal of Computational and Graphical Statistics*, 24(1), 44–65. <http://doi.org/10.1080/10618600.2014.907095>
- Goodman, B., & Flaxman, S. (2017). European Union Regulations on Algorithmic Decision-Making and a “Right to Explanation.” *AI Magazine*, 38(3), 50. <http://doi.org/10.1609/aimag.v38i3.2741>
- Gunning, D. (2017). Explainable artificial intelligence (xai). *Defense Advanced Research Projects Agency (DARPA), Nd Web*. article.

Bibliography (contd.)

- Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural networks. *Proceedings of the 34th International Conference on Machine Learning - Volume 70*. JMLR.org. Retrieved from <https://dl.acm.org/citation.cfm?id=3305518>
- He, J., Baxter, S. L., Xu, J., Xu, J., Zhou, X., & Zhang, K. (2019). The practical implementation of artificial intelligence technologies in medicine. *Nature Medicine*, 25(1), 30–36. <http://doi.org/10.1038/s41591-018-0307-0>
- Hosny, A., Parmar, C., Quackenbush, J., Schwartz, L. H., & Aerts, H. J. W. L. (2018). Artificial intelligence in radiology. *Nature Reviews Cancer*, 18(8), 500–510. <http://doi.org/10.1038/s41568-018-0016-5>
- Jiang, H., Kim, B., Guan, M., & Gupta, M. (2018). To trust or not to trust a classifier. In *Advances in Neural Information Processing Systems* (pp. 5546–5557). inproceedings.
- Jungo, A., Meier, R., Ermis, E., Herrmann, E., & Reyes, M. (2018). Uncertainty-driven Sanity Check: Application to Postoperative Brain Tumor Cavity Segmentation. In *1st Conference on Medical Imaging with Deep Learning (MIDL 2018)* (Vol. In Press). conference.
- Kim, B., Wattenberg, M., Gilmer, J., Cai, C., Wexler, J., Viegas, F., & Sayres, R. (2017). Interpretability Beyond Feature Attribution: Quantitative Testing with Concept Activation Vectors (TCAV). Retrieved from <http://arxiv.org/abs/1711.11279>
- Kindermans, P.-J., Schütt, K. T., Alber, M., Müller, K.-R., Erhan, D., Kim, B., & Dähne, S. (2017). Learning how to explain neural networks: PatternNet and PatternAttribution, 1–12. <http://doi.org/10.1021/jm900403j>
- Kleesiek, J., Petersen, J., Döring, M., Maier-Hein, K., Köthe, U., Wick, W., ... Biller, A. (2016). Virtual Raters for Reproducible and Objective Assessments in Radiology. *Scientific Reports*, 6, 25007. <http://doi.org/10.1038/srep25007>
- Koh, P. W., & Liang, P. (2017). Understanding Black-box Predictions via Influence Functions. <http://doi.org/10.3389/fpsyg.2012.00223>
- Lambin, P., Leijenaar, R. T. H., Deist, T. M., Peerlings, J., de Jong, E. E. C., van Timmeren, J., ... Walsh, S. (2017). Radiomics: the bridge between medical imaging and personalized medicine. *Nature Reviews Clinical Oncology*, 14(12), 749–762. <http://doi.org/10.1038/nrclinonc.2017.141>
- Lapuschkin, S., Wäldchen, S., Binder, A., Montavon, G., Samek, W., & Müller, K.-R. (2019). Unmasking Clever Hans Predictors and Assessing What Machines Really Learn. <http://doi.org/10.1038/s41467-019-08987-4>
- Letham, B., Rudin, C., McCormick, T. H., & Madigan, D. (2015). Interpretable classifiers using rules and Bayesian analysis: Building a better stroke prediction model. *Ann. Appl. Stat.*, 9(3), 1350–1371. article. <http://doi.org/10.1214/15-AOAS848>
- Lipton, Z. C. (2016). The Mythos of Model Interpretability. In *ICML Workshop on Human Interpretability in Machine Learning (WHI)*. Retrieved from <http://arxiv.org/abs/1606.03490>
- Litjens, G., Kooi, T., Bejnordi, B. E., Setio, A. A. A., Ciompi, F., Ghafoorian, M., ... Sánchez, C. I. (2017, December 1). A survey on deep learning in medical image analysis. *Medical Image Analysis*. Elsevier. <http://doi.org/10.1016/j.media.2017.07.005>
- Mahapatra, D., Bozorgtabar, B., Thiran, J., & Reyes, M. (2018). Efficient Active Learning for Image Classification and Segmentation using a Sample Selection and Conditional Generative Adversarial Network. In *Medical Image Computing and Computer-Assisted Intervention -- MICCAI 2018* (Vol. In Press). inproceedings.
- Maier-Hein, L., Ross, T., Gröhl, J., Glocker, B., Bodenstedt, S., Stock, C., ... Maier-Hein, K. (2016). Crowd-Algorithm Collaboration for Large-Scale Endoscopic Image Annotation with Confidence (pp. 616–623). Springer, Cham. http://doi.org/10.1007/978-3-319-46723-8_71
- Miller, T. (2017). Explanation in Artificial Intelligence: Insights from the Social Sciences. Retrieved from <http://arxiv.org/abs/1706.07269>

Bibliography (contd.)

- Moosavi-Dezfooli, S.-M., Fawzi, A., & Frossard, P. (2016). DeepFool: A Simple and Accurate Method to Fool Deep Neural Networks. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 2574–2582). IEEE. <http://doi.org/10.1109/CVPR.2016.282>
- Murdoch, W. J., Singh, C., Kumbier, K., Abbasi-Asl, R., & Yu, B. (2019). *Interpretable machine learning: definitions, methods, and applications*. Retrieved from <https://arxiv.org/pdf/1901.04592.pdf>
- Nair, T., Precup, D., Arnold, D. L., & Arbel, T. (2018). Exploring Uncertainty Measures in Deep Networks for Multiple Sclerosis Lesion Detection and Segmentation (pp. 655–663). Springer, Cham. http://doi.org/10.1007/978-3-030-00928-1_74
- Nguyen, A., Yosinski, J., & Clune, J. (2015). Deep neural networks are easily fooled: High confidence predictions for unrecognizable images. In *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 427–436). IEEE. <http://doi.org/10.1109/CVPR.2015.7298640>
- Parikh, R. B., Obermeyer, Z., & Navathe, A. S. (2019). Regulation of predictive analytics in medicine. *Science*, 363(6429), 810–812. <http://doi.org/10.1126/science.aaw0029>
- Pereira, S., Meier, R., McKinley, R., Wiest, R., Alves, V., Silva, C. A., & Reyes, M. (2018). Enhancing interpretability of automatically extracted machine learning features: application to a RBM-Random Forest system on brain lesion segmentation. *Medical Image Analysis*, 44, 228–244. <http://doi.org/10.1016/j.media.2017.12.009>
- Poursabzi-Sangdeh, F., Goldstein, D. G., Hofman, J. M., Vaughan, J. W., & Wallach, H. (2018). Manipulating and Measuring Model Interpretability. Retrieved from <http://arxiv.org/abs/1802.07810>
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). Why Should I Trust You?. Explaining the Predictions of Any Classifier. Retrieved from <http://arxiv.org/abs/1602.04938>
- Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., & Batra, D. (2017). Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization. In *2017 IEEE International Conference on Computer Vision (ICCV)* (pp. 618–626). IEEE. <http://doi.org/10.1109/ICCV.2017.74>
- Simonyan, K., Vedaldi, A., & Zisserman, A. (2013). Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps. Retrieved from <http://arxiv.org/abs/1312.6034>
- Szegedy, Christian, Zaremba, Wojciech, Sutskever, Ilya, Bruna, Joan, Erhan, Dumitru, Goodfellow, Ian, & Fergus, Rob. (2013). Intriguing properties of neural networks. *Eprint arXiv:1312.6199*. Retrieved from <http://adsabs.harvard.edu/abs/2013arXiv1312.6199S>
- Topol, E. J. (2019). High-performance medicine: the convergence of human and artificial intelligence. *Nature Medicine*, 25(1), 44–56. <http://doi.org/10.1038/s41591-018-0300-7>
- Van Lent, M., Fisher, W., & Mancuso, M. (2004). An explainable artificial intelligence system for small-unit tactical behavior. In *Proceedings of the National Conference on Artificial Intelligence* (pp. 900–907). inproceedings.
- Zech, J. R., Badgeley, M. A., Liu, M., Costa, A. B., Titano, J. J., & Oermann, E. K. (2018). Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: A cross-sectional study. *PLOS Medicine*, 15(11), e1002683. <http://doi.org/10.1371/journal.pmed.1002683>